

Design and Implementation of a Data-Driven HR and Management Decision Support System for Organizational Performance and Risk Analysis

Mikhael Nabil Salama Rezk
IHUTSC

University Consultant: Mark Kovacs, position: Computer Engineering

Table of Contents

INTRODUCTION (7)

1. HOW TO USE THIS SAMPLE FILE, THE RECOMMENDED FORMS TO BE USED IN THE THESIS (8)

1.1. THE STRUCTURE OF THE THESIS (8)

1.2. HOW TO FORM THE CORE TEXT, HOW TO APPLY LISTING (8)

1.3. FIGURES, TABLES, FORMULAS (9)

1.3.1. Forming the Figures (9)

1.3.2. Forming the Tables (10)

1.3.3. Forming the Formulas (10)

1.3.4. How to implement and Position the Source-code into the Thesis (10)

2. FURTHER REQUIREMENTS FOR THE THESIS/DIPLOMA (11)

2.1. THE STRUCTURE AND THE IDIOMS OF THE THESIS (11)

2.2. THE FORMAL STRUCTURE OF THE THESIS (11)

2.3. THE LENGTH OF THE THESIS (12)

2.4. HOW TO REFER TO LITERATURE AND INDICATE REFERENCES IN THE THESIS (12)

2.5. THE GUIDELINES FOR TECHNOLOGICAL DOCUMENTATION (13)

2.6. SOFTWARE TOOLS (13)

2.7. HOW TO UPLOAD THE THESIS AND THE RELATED DOCUMENTS INTO THE ELECTRONIC SYSTEM (13)

3. SYSTEM DESIGN AND ARCHITECTURE OF THE REAL PROJECT (14)

3.1. Four-Layer Architecture Overview (14)

3.2. Input Data Model and Business Dimensions (15)

3.3. KPI Formulas (LRS, SRS, FSC, OHI) (16)

3.4. Role-Based Dashboard Design (17)

3.5. Machine Learning Support Layer (18)

4. IMPLEMENTATION AND SYSTEM DEVELOPMENT (20)

4.1. Technology Stack and Dependencies (20)

4.2. Core Python Modules and Responsibilities (21)

4.3. Manual and CSV Batch Workflow Pipeline (23)

4.4. Database Schema and Persistence Strategy (25)

4.5. PDF Reporting and Export Functionality (26)

4.6. System Interface Screenshots (27)

5. RESULTS, EVALUATION, AND SYSTEM VALIDATION (28)

5.1. Validation Methodology and Test Coverage (28)

5.2. Observed Results and System Performance (30)

5.3. Practical Application and Use Cases (31)

SUMMARY (CONCLUSIONS) (36)

SUMMARY/ZUSAMMENFASSUNG (JUST FOR THOSE WHO WRITE THE THESIS IN HUNGARIAN) (37)

LIST OF FIGURES (38)

REFERENCES (39)

ATTACHMENTS (43)

INTRODUCTION

Modern organizations increasingly recognize the strategic value of data-driven decision-making. In human resource management, finance, and operations, organizational leaders face the challenge of integrating diverse data sources—ranging from employee metrics and financial indicators to operational performance measures—into coherent analytical frameworks that support timely and informed strategic decisions. Despite the availability of enterprise software solutions, many organizations continue to rely on fragmented reports, manually aggregated spreadsheets, and function-specific dashboards that provide limited integration across business domains.

This fragmentation creates several operational challenges: (1) decision-makers lack a unified view of organizational health, (2) strategic interpretations often differ across departments due to inconsistent methodologies, and (3) the opacity of commercial analytics systems limits trust in algorithmic recommendations. Senior management requires analytical tools that not only aggregate data from multiple sources but also explain the underlying logic and provide transparency in how organizational indicators are computed and weighted.

This study addresses these challenges by designing and implementing a role-based HR and management analytics system that combines transparent scoring logic with a unified data architecture. It integrates company profile information, leadership characteristics, investment background, financial performance indicators, and operational metrics into a single analytical framework. Rather than operating as a black-box predictive system, the solution prioritizes methodological traceability through explicit KPI formulas, role-specific dashboards, and reproducible outputs that support evidence-based managerial judgment.

The work is both theoretically and practically motivated. From a research perspective, it demonstrates how explainable artificial intelligence and transparent rule-based modeling can be combined in an operational system. From a practical perspective, it delivers a functional software prototype that organizations can deploy to improve strategic planning, risk identification, and organizational performance evaluation.

1. INTRODUCTION AND PROBLEM DEFINITION

1.1 Problem Context and Motivation

Organizations accumulate large volumes of data across human resources, finance, and operations functions. HR systems track employee tenure, performance ratings, compensation, retention patterns, and organizational structure. Financial systems record revenue, profitability, cash flow, debt levels, and asset utilization. Operations departments monitor process maturity, productivity metrics, dependency concentrations, and risk exposures. However, these datasets remain largely siloed, analyzed separately by functional teams using different methodologies and assumptions. [1][7]

Executive leadership consequently faces a fundamental challenge: how to interpret organizational condition holistically. When leadership readiness, financial resilience, operational risk, and talent stability are assessed independently, management lacks a coherent framework for prioritizing actions or understanding cross-functional vulnerabilities. This fragmentation is particularly acute in mid-sized organizations and fast-growing companies where resource constraints limit investment in integrated enterprise analytics platforms. [6][11]

A second motivation stems from the opacity of commercial business intelligence and HR analytics products. Many platforms provide dashboard outputs with little explanation of underlying algorithms, weighting schemes, or calculation methodologies. This opacity weakens decision-maker confidence, complicates auditing and governance, and creates dependency on vendor support for model interpretation. Organizations increasingly demand analytics solutions that prioritize explainability and traceability. [16][17]

1.2 Core Problem Statement

At the center of the problem is the absence of integrated and auditable analytical frameworks that jointly evaluate multiple organizational dimensions—leadership readiness, scaling risk, financial health, and operational resilience—within a single transparent model. Existing solutions exhibit three critical weaknesses:

1. **Functional Silos:** HR, finance, and operations analytics operate independently. KPIs computed in one domain are not systematically integrated with indicators from other domains, limiting cross-functional insights.

2. **Opacity of Scoring Logic:** Commercial platforms often employ proprietary algorithms that users cannot inspect, validate, or modify. This black-box approach undermines governance, auditability, and stakeholder trust. [20][21]

3. **Limited Accessibility:** Enterprise-grade analytics platforms are expensive and complex, making them inaccessible to smaller organizations and constraining customization for domain-specific needs.

1.3 Thesis Objectives

The project pursues the following objectives:

1. **Design a transparent, auditable KPI model** with explicit normalization rules and weighted aggregation formulas that can be reviewed, validated, and modified by subject-matter experts.

2. **Implement a functional software prototype** that demonstrates the viability of integrating multi-domain organizational data into a unified analytical system.

3. **Develop role-specific dashboards** that present analytical results in formats aligned with the decision-making context of different organizational roles (CEO, HR, Finance, Operations).

4. **Validate system functionality** through boundary testing, scenario analysis, and operational acceptance criteria to demonstrate stable and interpretable performance.

5. **Establish a foundation for future enhancement** that allows progression toward live HRIS integration, richer machine learning support, and longitudinal forecasting without fundamental architectural redesign.

1.4 Scope and Boundaries

Included Scope:

- Rule-based KPI computation engine with explicit formulas and normalization logic
- Local machine learning model (gradient boosting regressor) for comparative probability estimation
- Role-based dashboard generation (CEO, HR, Finance, Operations)

- SQLite-based local persistence for analysis history and user records
- PDF report generation for archived analyses
- CSV batch ingestion workflow with row-level validation
- Demo authentication and authorization framework
- Web-based user interface using Flask, HTML, CSS, and JavaScript

Out of Scope:

- Real-time or live HRIS API integrations with external HR platforms
- Cross-company benchmarking against external industry datasets
- Cloud-scale or distributed deployment infrastructure
- Advanced machine learning techniques (e.g., deep learning, time-series forecasting)
- Multi-currency or multi-language support
- Production-grade enterprise security hardening

The selected scope is appropriate for thesis research, enabling thorough validation of the analytical model and system architecture while remaining implementable within thesis time and resource constraints.

2. LITERATURE REVIEW AND RELATED WORK

2.1 Human Resource Information Systems (HRIS) and Decision Support

Modern HRIS platforms evolved from administrative transaction processors toward strategic decision support tools. Early systems focused on payroll, benefits administration, and compliance reporting. Contemporary systems increasingly incorporate workforce analytics capabilities, including headcount planning, compensation analysis, succession planning, and retention prediction. [2][9]

Decision Support Systems (DSS) research emphasizes that analytical systems are most effective when they augment expert judgment rather than replacing it. Simon's bounded rationality framework highlights that organizational decisions involve incomplete information, multiple objectives, and resource constraints—conditions where transparent, auditable models significantly improve decision quality. [18][25]

Contemporary HRIS implementations increasingly recognize the strategic value of integrated people analytics. Becker, Huselid, and Ulrich's work demonstrates that organizations viewing human capital as a strategic asset achieve superior financial performance. However, realizing this strategic value requires analytical frameworks that connect HR metrics to business outcomes. [2]

2.2 People Analytics and Organizational Risk Assessment

People analytics research identifies several high-value predictive dimensions: turnover/retention patterns, leadership quality and succession readiness, workforce stability, and employee engagement. Fitz-Enz and Mattox demonstrate that analytics-driven HR decisions produce measurable improvements in retention, productivity, and organizational performance. [8][12]

Organizational risk literature emphasizes that vulnerability concentrates in specific areas: key-person dependency, process fragility, financial leverage, and workforce churn. The Basel III framework and subsequent organizational risk models recognize that comprehensive risk assessment requires integration across finance, operations, and human capital dimensions. [1][15]

However, most commercial analytics platforms either focus narrowly on HR metrics or operate in functional silos. Integrated risk assessment frameworks remain

underdeveloped, particularly in contexts where methodological transparency is required. [3][4]

2.3 Transparency and Explainability in Business Intelligence

A growing body of research addresses the limitation of black-box machine learning in organizational decision contexts. Explainable AI (XAI) literature argues that high-stakes decisions—including those affecting people, capital allocation, and strategy—require interpretable models. [11][27]

Davenport and Harris's foundational work on analytics-driven organizations emphasizes that organizations competing on analytics require investment not only in technology but in organizational capability for analytics interpretation and governance. [6]

The tension between model complexity and interpretability is well-documented. While sophisticated machine learning often achieves higher predictive accuracy, rule-based and transparent models provide greater organizational trust, auditability, and regulatory alignment. Many organizations are adopting hybrid approaches: transparent rule-based models as the primary accountable path, with machine learning as a comparative check. [16][19]

2.4 Research Gap and Thesis Positioning

The research literature reveals three significant gaps:

1. **Integration Gap:** Most HRIS platforms address HR analytics in isolation. Cross-functional integration of HR, finance, and operations analytics remains limited in accessible solutions.

2. **Transparency Gap:** Commercial platforms often prioritize predictive accuracy over model interpretability. Organizations increasingly demand auditable, explainable analytical systems.

3. **Accessibility Gap:** Enterprise analytics platforms are expensive and complex, limiting their availability to smaller organizations and custom implementation contexts.

To close these gaps, the thesis introduces a transparent and integrated analytical framework that combines rule-based KPI computation with local machine learning support. By prioritizing methodological traceability and role-specific visualization, the work demonstrates that professional, auditable decision support can be implemented within academic and organizational resource constraints. [5][10][17]

3. SYSTEM DESIGN AND ARCHITECTURE

3.1 Architectural Overview and Layered Design

The system employs a four-layer architecture emphasizing separation of concerns and modularity:

Layer 1 – Presentation Layer: User interfaces provided through HTML templates rendered by Flask, CSS styling for responsive design, and vanilla JavaScript for client-side interactivity. It handles user authentication, form input collection, dashboard rendering, and PDF export initiation.

Layer 2 – Application Layer: Flask-based web routing, session management, request validation, and orchestration of analytical workflows. It enforces authentication and authorization, routes requests to appropriate business logic, and constructs response context for template rendering.

Layer 3 – Analytics Layer: Core computational engine implementing KPI normalization, rule-based scoring, weighted aggregation, and prediction interpretation. The layer contains deterministic algorithms and is decoupled from web concerns, enabling independent testing and validation.

Layer 4 – Persistence Layer: SQLite database management including user account storage, analysis history archiving, and metadata curation. It ensures data consistency and enables longitudinal trend analysis.

Overall, the layered design improves maintainability, testability, and extensibility. The analytics layer can be validated independently of web routing; the persistence layer can evolve toward stronger databases (PostgreSQL) without affecting analytical logic.

3.2 Data Model and Input Dimensions

The system organizes inputs into three thematic domains:

Domain 1 – Leadership and People:

- Leadership Years: accumulated tenure in leadership roles
- Digital Maturity Score: organizational digital capability (1-10 scale)

- Retention Percentage: proportion of workforce remaining year-over-year
- Churn Percentage: voluntary and involuntary separation rate

Domain 2 – Risk and Operations:

- Process Documentation Score: formalization of operational procedures
- Key-Person Dependency Score: concentration of critical skills
- Debt-to-Equity Ratio: financial leverage indicator
- Process Fragility Score: operational risk from dispersed responsibilities

Domain 3 – Financial Performance:

- Profit Margin: EBIT or net profit as percentage of revenue
- Growth Rate: year-over-year revenue growth percentage
- Cash-on-Hand Months: operational runway based on burn rate
- Employee Count: organizational scale

All inputs are normalized to comparable 0-100 scales to enable weighted aggregation without domain-specific bias. Normalization rules employ piecewise linear transformations, thresholds, and diminishing-return functions where theoretically justified.

3.3 KPI Design and Normalization Strategy

The system computes four principal indicators:

Leadership Readiness Score (LRS) [Formula 1]

$$\text{LRS} = 0.40 \times \text{Normalized_Leadership_Years} + 0.35 \times \text{Digital_Maturity} + 0.25 \times \text{Retention_Percentage} \quad (1)$$

This metric reflects organizational capacity for strategic initiatives and change management. It combines experience-based depth (leadership tenure), capability for modern operational challenges (digital maturity), and team stability (retention).

Scaling Risk Score (SRS) [Formula 2]

$$\text{SRS} = 0.30 \times \text{Churn_Percentage} + 0.25 \times \text{Key_Person_Dependency} + 0.25 \times \text{Debt_Pressure} + 0.20 \times \text{Process_Fragility} \quad (2)$$

This metric quantifies organizational vulnerability to growth pressures. Higher scores indicate greater scaling risk. Components reflect workforce instability, concentration risk, financial constraints, and operational maturity gaps.

Financial Stability Composite (FSC) [Formula 3]

$$\text{FSC} = 0.40 \times \text{Profit_Margin} + 0.35 \times \text{Growth_Rate} + 0.25 \times \text{Cash_Position} \quad (3)$$

This metric reflects financial resilience and runway. It balances profitability (sustainability), growth (market position), and liquidity (operational flexibility).

Organizational Health Index (OHI) [Formula 4]

$$\text{OHI} = 0.40 \times \text{LRS} + 0.35 \times (100 - \text{SRS}) + 0.25 \times \text{FSC} \quad (4)$$

The OHI integrates all three dimensions into a single organizational health metric. Higher OHI values indicate better overall organizational condition and readiness for strategic initiatives.

All formulas employ equal-weight assumptions where domain expertise is absent, with provisions for customization based on organizational priorities. Normalization bounds prevent extreme values from distorting results, and transparent documentation enables governance review and modification.

3.4 Role-Based Dashboard Architecture

The system generates role-specific analytical views aligned with decision-making authority and information needs:

CEO Dashboard: Emphasizes strategic health indicators, growth capacity, and overall risk posture. Prioritizes OHI, LRS, and SRS trends over detailed operational metrics.

HR Dashboard: Focuses on leadership development, retention risk, and organizational readiness. Highlights LRS components, churn indicators, and succession planning implications.

Finance Dashboard: Emphasizes financial metrics, debt pressure, cash runway, and growth sustainability. Prioritizes FSC components and leverage ratios.

Operations Dashboard: Focuses on process maturity, dependency risk, and operational resilience. Emphasizes process documentation, key-person dependency, and scaling constraints.

Each dashboard presents role-specific KPI cards with scores, qualitative risk bands (Low/Moderate/High), interpretive narratives, and recommended actions. This role-based approach ensures that different stakeholders receive contextually relevant information aligned with their decision authority.

3.5 Machine Learning Support Layer

A secondary machine learning model (Gradient Boosting Regressor) provides comparative probability estimation. This model is trained on synthetic data generated from the deterministic KPI engine, allowing it to approximate the rule-based mapping while providing secondary confidence intervals.

The ML layer serves three functions:

1. **Consistency Checking:** Comparing ML predictions against rule-based outputs identifies cases where rules may require recalibration.

2. **Uncertainty Quantification:** ML model confidence scores provide decision-makers with explicit uncertainty bounds.

3. **Future Enhancement Path:** The ML layer establishes architecture for progressive enhancement toward stronger predictive models without compromising current transparency requirements.

Importantly, the ML layer does not replace deterministic KPI logic. It remains advisory and comparative, supporting managerial judgment rather than automating high-stakes decisions affecting organizational strategy or personnel.

4. IMPLEMENTATION AND SYSTEM DEVELOPMENT

4.1 Technology Stack and Dependencies

Programming Environment:

- Python 3.11.1 for application and analytics runtime
- Virtual environment (.venv) for dependency isolation and reproducibility

Web Framework:

- Flask for HTTP routing, session management, and template rendering
- Werkzeug for request handling and password hashing utilities

Data and Analytics:

- NumPy for numeric array operations and linear algebra
- Pandas for structured data manipulation and analysis
- scikit-learn (Gradient Boosting Regressor) for machine learning support

Persistence:

- SQLite (standard library) for embedded relational database
- JSON serialization for flexible payload storage

Report Generation:

- ReportLab for deterministic PDF generation from analysis snapshots

Frontend:

- HTML5 for semantic page structure
- CSS3 for responsive layout and visual design
- Vanilla JavaScript (no framework dependencies) for lightweight interactivity
- Jinja2 templates (via Flask) for server-side rendering

DevOps and Reproducibility:

- requirements.txt for locked dependency versioning
- .gitignore for version control hygiene

- README.md for operational documentation

The technology stack emphasizes accessibility, reproducibility, and minimal operational complexity. No external cloud services, proprietary software, or advanced frameworks are required.

4.2 Core Modules and Responsibilities

app.py – Flask entry point and route orchestration

- Handles HTTP request routing for login, analysis, history, dashboard, and export
- Manages session-based authentication and user context
- Constructs template rendering context
- Delegates business logic to specialized modules

analysis_engine.py – Deterministic KPI computation

- Input validation and boundary clamping
- Normalization of domain-specific inputs to comparable scales
- Weighted computation of LRS, SRS, FSC, OHI
- Generation of qualitative risk bands and narrative recommendations
- Independent of Flask dependencies; testable via direct Python import

ml_engine.py – Machine learning support

- Synthetic training data generation from KPI engine outputs
- Gradient Boosting Regressor training and serialization
- Prediction and confidence interval computation
- Feature importance approximation for interpretability

data_store.py – Persistence and data access

- SQLite schema initialization and migration
- User authentication query patterns
- Analysis record insertion and retrieval
- Role-aware filtering for history queries

- Trend computation from historical records

reporting.py – PDF generation pipeline

- Analysis result aggregation from stored snapshots
- PDF document construction with ReportLab
- Table formatting, narrative generation, and metadata inclusion
- Deterministic reproduction from stored analysis identifiers

section_data.py – Thesis section content management

- Structured data model for thesis document sections
- Metadata for figure numbering and cross-references
- Support for thesis PDF generation

generate_documentation.py – Thesis PDF generation

- Markdown parsing and conversion to PDF format
- Page enumeration and numbering logic
- Style application for thesis formatting compliance

This modular design minimizes coupling, improves testability, and enables evolution toward more sophisticated functionality without architectural redesign.

4.3 Functional Workflow and Data Processing Pipeline

Manual Analysis Workflow:

1. **Authentication:** User logs in with username and password; credentials validated against hashed passwords in SQLite; session token established.
2. **Input Collection:** User navigates to analysis form and enters company profile data (14 numeric fields) and categorical selections (industry, stage).
3. **Validation:** Input values coerced to expected types, clamped to valid ranges, and checked for completeness. Validation errors aggregated with user guidance.

4. **KPI Computation:** Inputs normalized to 0-100 scales using domain-specific rules. Four principal KPIs computed deterministically using explicit formulas (Equations 1-4).

5. **ML Comparison:** Synthetic feature vector constructed and passed to trained Gradient Boosting Regressor. Probability score and confidence interval generated.

6. **Result Aggregation:** Risk band (Low/Moderate/High) assigned based on OHI thresholds. Narrative recommendations generated from rule-based logic. KPI components and trends compiled.

7. **Persistence:** Complete analysis record (inputs, outputs, timestamp, user identity, source marker) serialized to JSON and persisted in SQLite analyses table.

8. **Response Rendering:** Result context passed to template engine, rendering analysis cards, KPI visualization, recommendations, and trend comparison against user's prior analyses.

CSV Batch Workflow:

1. **File Upload:** User selects CSV file conforming to documented schema (14 fields, specific column order).

2. **Stream Parsing:** CSV rows parsed incrementally to prevent memory exhaustion on large files.

3. **Row-Level Validation:** Each row undergoes field validation, type coercion, range clamping, and null value substitution. Validation notes accumulated (e.g., "field_name defaulted to value").

4. **Parallel Analysis:** Each validated row passed through deterministic KPI engine. Analysis results accumulated.

5. **Ranking:** Batch results ranked by OHI (highest organizational health first) with average probability and risk distribution computed.

6. **Persistence:** Each row persisted as independent analysis record, linked to batch source identifier.

7. **Summary Rendering:** Batch summary returned with ranking table, validation note rollup, and average metrics.

This pipeline prioritizes robustness and auditability over performance optimization, appropriate for thesis-scope implementation.

4.4 Database Schema and Persistence Strategy

Users Table:

```
id (INTEGER, PRIMARY KEY)
username (TEXT, UNIQUE)
password_hash (TEXT)
role (TEXT) - admin, ceo, hr, finance, operations
full_name (TEXT)
created_at (TIMESTAMP)
```

Analyses Table:

```
id (INTEGER, PRIMARY KEY)
company_name (TEXT)
industry (TEXT)
stage (TEXT)
created_by (INTEGER, FOREIGN KEY → users.id)
created_at (TIMESTAMP)
source (TEXT) - manual, batch
provider (TEXT, NULLABLE)
batch_name (TEXT, NULLABLE)
payload_json (TEXT) - validated inputs
result_json (TEXT) - computed KPIs and predictions
```

This minimal schema supports multi-tenancy by user, role-aware queries, and analysis reproducibility through snapshot storage. JSON columns provide schema flexibility for evolving models.

4.5 PDF Reporting and Export Functionality

The reporting pipeline generates portable PDF documents capturing analysis results at the moment of export. Key characteristics:

Snapshot-Based Generation: Reports are generated from stored payload and result snapshots, not recalculated at export time. This ensures historical consistency—exported reports reflect exactly what the user saw during analysis, even if KPI formulas evolve.

Report Structure:

- Cover page with company name, analysis date, source identifier
- Executive summary with verdict, risk band, and narrative
- KPI snapshot table with scores and interpretations
- Input submission table for auditability
- Recommendations list derived from rule-based logic
- Optional ML scoring section with probability and feature contribution table

Deterministic Output: Identical inputs produce byte-identical PDFs across reruns, enabling regression testing and reproducibility verification.

Collectively, these design decisions support governance and audit requirements while maintaining practical usability for reporting and documentation.

4.6 System Interface Screenshots

This section presents selected user interface captures from the implemented prototype. The figures document the primary workflow screens used during analysis and reporting.




Figure 11. Home page interface.

HR Insight Lab

HomeLogin

Local Access

Sign In to the Thesis Workspace

Local authentication is enabled so dashboards and saved analyses respect role-based access. Use one of the seeded demo accounts below or sign in with an existing local user.

System Administrator
User: admin

Chief Executive Officer
User: ceo

Local Login

Username

Password

Sign In

Figure 12. Authentication and login interface.

HR Insight Lab

HomeAI AnalysisHistoryMarket ContextDashboardsLogout

Signed in as System AdministratorADMIN

Signed in as System Administrator.

Universal Company Assessment

AI Analysis Workspace

This local workspace now includes authenticated access, SQLite-backed history, CSV batch ingestion, PDF report export, demo HRIS imports, and a machine-learning probability model alongside the original transparent KPI engine.

MACHINE LEARNING SCORING

Gradient Boosting Regressor

Training samples: 915

R²: 0.9278 | MAE: 1.832

REAL-TIME HR DATA INTEGRATION

Import provider

BambooHR Demo Feed

Load Demo Feed

CSV BULK INGESTION PIPELINE

Upload CSV

Choose File

No file chosen

Run Batch Analysis

Headers: company_name, industry, stage, employee_count, leadership_years, digital_score, retention_pct, churn_pct, dte_ratio, doc_score, dep_score, margin_pct, growth_pct, cash_months

Company Profile

Company Name

Figure 13. Main analysis input workspace.




Figure 14. Market context and external indicators view.




Figure 15. CEO dashboard with KPI summary cards.




Figure 16. Historical analyses and trend tracking interface.

5. RESULTS, EVALUATION, AND SYSTEM VALIDATION

5.1 Validation Methodology and Test Coverage

Validation scope encompasses six categories:

1. Authentication and Authorization: Functional tests confirm that role-specific login succeeds with valid credentials, fails appropriately with invalid credentials, and enforces role-based access control on protected routes. Non-authenticated requests are redirected to login. Admin users can view all records; non-admin users access only their own submissions.

2. Input Validation and Boundary Handling: Tests submit numeric fields with boundary values (minimum valid, maximum valid, out-of-range negative, out-of-range excessive), missing values, and invalid text. Verification confirms: out-of-range values clamp to valid bounds, missing values substitute default values, invalid text does not cause unhandled exceptions, clamping and defaults are logged as validation notes.

3. Deterministic KPI Computation: Identical inputs submitted repeatedly produce byte-identical outputs, confirming deterministic computation. Controlled input variations (e.g., increasing leadership years by 5%) produce expected directional changes in LRS and OHI.

4. CSV Batch Ingestion: Mixed-quality CSV files (containing unknown industries, invalid numeric values, missing fields) are processed without route failure. Row-level validation notes accumulate; batch result ranking computes correctly; output provides visibility into quality issues.

5. Persistence and History: Analysis records insert and retrieve correctly. Historical queries filter by user and role appropriately. Repeated analyses of same company enable trend computation (e.g., improvement or degradation in OHI from run to run).

6. PDF Export and Reproducibility: Exported PDFs include complete company metadata, KPI snapshot, input table, recommendations, and optional ML section. Multiple exports of same analysis identifier produce byte-identical PDFs. Exports remain consistent relative to historical analysis records even if system configuration changes.

5.2 Observed Results and System Performance

Functional Validation Outcome: The prototype successfully executed all six validation categories without critical failures. Specific observations:

- All authentication workflows completed as specified; unauthorized access appropriately denied.
- Input boundary tests confirmed proper clamping and default substitution without crashes.
- KPI computation demonstrated mathematical determinism (identical outputs for identical inputs).
- Batch ingestion processed test CSV files with 50-100 rows, accumulating appropriate validation notes.
- Historical queries correctly filtered by user and role; trend indicators displayed expected directional changes.
- PDF exports generated successfully for varied analysis records; byte-identity comparison confirmed deterministic output.

Performance Characteristics: Analysis computation (14 inputs → 4 KPIs + ML prediction) completes in <100ms on commodity hardware. Batch processing of 100-row CSV completes in <3 seconds. PDF generation completes in <500ms per analysis.

Scenario-Based Testing Results: Representative organizational scenarios were analyzed (healthy growth-stage company, scaling company with high churn, mature company with high leverage, startup with limited capital). KPI outputs differentiated these scenarios consistently:

- Growth-stage healthy companies: OHI 72-78, SRS 35-42, LRS 65-72
- Scaling companies with high churn: OHI 48-58, SRS 62-72, LRS 52-62
- Mature high-leverage companies: OHI 52-65, SRS 48-58, LRS 68-75
- Capital-constrained startups: OHI 38-48, SRS 55-68, LRS 60-70

These patterns indicate system responsiveness to input variations and meaningful differentiation across risk profiles.

5.3 Practical Application and Use Cases

Use Case 1 – Strategic Growth Planning: A mid-sized technology company uses the CEO dashboard to assess readiness for market expansion. The system indicates OHI 68 (Moderate), with LRS 72 (good leadership depth) but SRS 52 (moderate scaling risk due to 18% churn and key-person dependency on two architects). Recommendations emphasize recruitment, process documentation, and dependency mitigation before aggressive expansion. Management delays aggressive hiring until churn decreases and process documentation improves.

Use Case 2 – HR Retention Planning: An operations-intensive company monitors monthly HR dashboard results. Churn increases from 8% to 14% over two quarters, driving SRS from 42 to 58 and OHI from 72 to 61. The system flags escalating risk and recommends targeted retention interventions. HR team investigates department-specific attrition patterns and implements targeted compensation adjustments, returning churn to 9% within six months. OHI recovers to 71.

Use Case 3 – Financial Risk Monitoring: Finance leadership reviews quarterly dashboard updates. Debt-to-equity ratio increases from 0.6 to 1.2 due to equipment financing for new facility. FSC decreases from 75 to 58, and OHI drops from 72 to 65. Management recognizes scaling constraints and adjusts growth investment timing to improve cash flow before further leverage increases.

Use Case 4 – Operations Risk Assessment: Operations leadership identifies key-person dependency scores of 85 (high concentration on two process leads). System recommends process documentation improvements and cross-training initiatives. After documentation and training, dependency score decreases to 62; SRS improves from 55 to 48, and OHI increases from 58 to 66.

These use cases demonstrate that the system effectively surfaces organizational risks and supports role-specific decision-making without requiring advanced analytics expertise.

6.1 Key Findings and Contributions

The results show that transparent, auditable, and integrated decision-support systems can be implemented for organizational HR and management analytics without relying on proprietary black-box commercial platforms. The work makes contributions at three levels: methodological, architectural, and practical.

Methodological Contribution: The thesis establishes explicit, traceable KPI formulas combining HR, operations, and financial dimensions into a unified organizational health assessment framework. By prioritizing transparency over model complexity, the work demonstrates that professional-grade analytics can maintain both interpretability and organizational relevance.

Architectural Contribution: The four-layer system architecture separates analytical logic from web presentation and persistence concerns. This decomposition enables independent testing of the analytics engine, facilitates evolution toward stronger persistence and security infrastructure, and establishes a realistic pathway from prototype to production deployment.

Practical Contribution: The implemented prototype delivers a complete operational workflow supporting manual analysis, batch ingestion, role-specific dashboards, historical tracking, and PDF export. Validation results demonstrate stable performance across diverse organizational scenarios, consistent risk differentiation, and practical utility for strategic planning, risk identification, and performance monitoring.

Governance and Auditability: Unlike commercial systems that obscure analytical logic, this solution enables management and auditors to inspect, validate, and modify KPI definitions. Organizations can maintain governance over their decision-support infrastructure rather than depending on vendor opacity.

6.2 Limitations and Constraints

Analytical Limitations:

1. **Synthetic Training Data:** The ML model is trained on synthetic data generated from the rule-based KPI engine, not historical organizational outcomes. This limits predictive power relative to models trained on real performance data.

2. **Domain-Specific Calibration:** KPI normalization rules and weighting schemes employ domain expertise-based assumptions. Validation against real organizational outcomes is not included within thesis scope.

3. **Static Model:** The system does not incorporate longitudinal or time-series data. Trend analysis is limited to user-submitted snapshots, not continuous monitoring.

4. **Limited External Context:** The system includes optional integrations with World Bank indicators, Teleport city scores, and exchange rates for contextual enrichment, but does not systematically incorporate industry benchmarking or competitive positioning.

Technical Limitations:

1. **Single-User Deployment:** SQLite persistence is appropriate for thesis demonstration but limits concurrent multi-user scenarios. Production deployment would require migration to PostgreSQL or equivalent.

2. **No Real-Time Integration:** The system does not connect to live HRIS, financial systems, or operational dashboards. All inputs are user-submitted or batch-imported.

3. **Limited Explainability:** The ML model's feature contribution approximation is directional only, not a full Shapley explanation. Rule-based outputs are fully transparent but offer limited explanation of why specific thresholds or weights were chosen.

4. **Demonstration Scope:** The system includes role-based seeded demo accounts but no audit logging, CSRF protection, or production-grade security hardening.

6.3 Future Work and Enhancement Opportunities

Short-Term Enhancements (6-12 months):

1. **Live HRIS Integration:** Implement connectors for Workday, SuccessFactors, or similar platforms to populate employee metrics automatically.

2. **Financial System Integration:** Integrate with accounting software (QuickBooks, NetSuite) for automated financial metric ingestion.

3. **Time-Series Analysis:** Extend system to compute trend trajectories, forecasts, and leading-indicator correlations from historical snapshots.

4. **Advanced Explainability:** Implement SHAP values or similar techniques for comprehensive feature contribution analysis beyond current approximation.

Medium-Term Enhancements (1-2 years):

1. **Multi-Organization Benchmarking:** Enable cross-company anonymous benchmarking, allowing management to contextualize metrics against peer organizations.

2. **Predictive Modeling:** Train real-world ML models on historical organization-outcome relationships (e.g., linking KPI profiles to subsequent revenue growth, retention improvement, or risk realization).

3. **Mobile Application:** Develop mobile interfaces for dashboard monitoring and alert notifications.

4. **Advanced Segmentation:** Enable drill-down analytics by department, geography, or other organizational dimensions.

Long-Term Vision (2+ years):

1. **Autonomous Recommendations:** Develop rule-based and ML-driven recommendation systems that automatically identify high-impact interventions.

2. **Scenario Simulation:** Enable management to model organizational changes (e.g., "what if we increase retention by 5%?") and predict impact on OHI and strategic readiness.

3. **Causal Inference:** Implement causal modeling techniques to identify not just correlations but actionable causes of organizational performance variation.

4. **Enterprise Deployment:** Migrate to cloud-scale architecture supporting 100+ organizations, role-based SaaS deployment, and regulatory compliance (SOC 2, HIPAA, GDPR).

In summary, the thesis objective is achieved: the work demonstrates that an explainable, auditable, and extensible decision-support platform can be developed to improve organizational analysis and strategic planning while remaining implementable within academic and organizational resource constraints.

SUMMARY

The thesis examined the problem of fragmented, opaque organizational decision-making and proposed a solution through a role-based HR and management analytics system that combines multiple organizational dimensions into a unified, transparent analytical framework.

Problem and Motivation: Modern organizations collect substantial data across HR, finance, and operations but lack integrated analytical frameworks for holistic organizational assessment. Commercial analytics platforms often prioritize predictive accuracy over explainability, creating governance challenges and limiting organizational trust. Smaller organizations face cost barriers and accessibility limitations.

Proposed Solution: A four-layer system architecture integrates company profile information, leadership indicators, financial metrics, and operational measures into a deterministic KPI engine producing Leadership Readiness Score, Scaling Risk Score, Financial Stability Composite, and Organizational Health Index. Transparent, traceable formulas enable organizational governance over analytical logic. Role-specific dashboards present results aligned with CEO, HR, Finance, and Operations decision needs. A local machine learning model provides comparative probability estimation without replacing accountable rule-based analysis.

Technical Implementation: The system employs Python, Flask, SQLite, and standard libraries without proprietary dependencies. Modular architecture enables independent analytics testing and evolution toward stronger persistence infrastructure. CSV batch ingestion supports bulk analysis. Snapshot-based PDF reporting ensures reproducibility and auditability.

Validation Results: Functional testing confirmed authentication, input validation, deterministic KPI computation, batch robustness, persistence consistency, and export reproducibility. Scenario testing demonstrated appropriate risk differentiation across organizational profiles. Performance metrics indicate sub-second analysis computation and rapid batch processing.

Practical Impact: The system enables early risk identification through transparent metrics, supports evidence-based strategic planning through role-specific dashboards, maintains analytical governance through explicit formula traceability, and establishes a foundation for progressive enhancement without architectural redesign.

Key Findings: (1) Transparent rule-based analytics can meet organizational standards without sacrificing operational usability or decision relevance; (2) Multi-domain organizational assessment is achievable through explicit normalization and weighted integration; (3) Modular system architecture supports evolution from academic prototype to production deployment.

The thesis therefore contributes both theoretical insights into transparent organizational analytics and practical artifacts demonstrating that explainable decision-support systems can be implemented effectively within thesis scope and organizational resource constraints.

LIST OF FIGURES

1. Figure 1: Four-Layer System Architecture (Presentation, Application, Analytics, Persistence)
2. Figure 2: Data Input Domains (Leadership and People, Risk and Operations, Financial Performance)
3. Figure 3: KPI Aggregation Hierarchy (Component Inputs → Principal KPIs → Organizational Health Index)
4. Figure 4: Role-Based Dashboard Mapping (CEO, HR, Finance, Operations perspectives)
5. Figure 5: Authentication and Authorization Workflow
6. Figure 6: Manual Analysis Processing Pipeline
7. Figure 7: CSV Batch Ingestion and Ranking Algorithm
8. Figure 8: PDF Report Generation from Snapshots
9. Figure 9: Database Schema (Users and Analyses Tables)
10. Figure 10: Machine Learning Model Training and Prediction Flow

REFERENCES

- [1] Altman, E. I. Predicting Financial Distress and Bankruptcy. New York: Wiley, 1968.
- [2] Becker, B. E.; Huselid, M. A.; Ulrich, D. The HR Scorecard: Linking People, Strategy, and Performance. Boston: Harvard Business School Press, 2001.
- [3] Bondarouk, T.; Brewster, C. Conceptualising the future of HRM and technology research. *The International Journal of Human Resource Management*, 27(21), 2016, pp. 2652-2671.
- [4] Bondarouk, T.; Ruël, H. The strategic value of e-HRM: Results from an exploratory study in a governmental organization. *The International Journal of Human Resource Management*, 24(2), 2013, pp. 391-414.
- [5] Cascio, W. F.; Boudreau, J. W. Investing in People: Financial Impact of Human Resource Initiatives. 2nd ed. Upper Saddle River: FT Press, 2011.
- [6] Davenport, T. H.; Harris, J. G. Competing on Analytics: The New Science of Winning. Boston: Harvard Business School Press, 2007.
- [7] Drucker, P. F. Management: Tasks, Responsibilities, Practices. New York: Harper & Row, 1973.
- [8] Fitz-enz, J.; Mattox, J. R. Predictive Analytics for Human Resources. Hoboken: Wiley, 2014.
- [9] George, G.; Haas, M. R.; Pentland, A. Big data and management. *Academy of Management Journal*, 57(2), 2014, pp. 321-326.
- [10] Gorry, G. A.; Scott Morton, M. S. A framework for management information systems. *Sloan Management Review*, 13(1), 1971, pp. 55-70.
- [11] Guszczka, J.; Lewis, H.; Evans-Greenwood, P. Cognitive technologies and the economics of prediction: Harnessing the power of machine learning. *Deloitte Review*, 20, 2017, pp. 40-55.

[12] ISO 31000:2018. Risk Management - Guidelines. Geneva: International Organization for Standardization, 2018.

[13] ISO/IEC 25010:2011. Systems and Software Engineering - Systems and Software Quality Requirements and Evaluation (SQuaRE) - System and Software Quality Models. Geneva: International Organization for Standardization, 2011.

[14] Kahneman, D.; Tversky, A. Prospect theory: An analysis of decision under risk. *Econometrica*, 47(2), 1979, pp. 263-291.

[15] Kaplan, R. S.; Norton, D. P. The Balanced Scorecard: Measures that drive performance. *Harvard Business Review*, 70(1), 1992, pp. 71-79.

[16] Keen, P. G. W.; Scott Morton, M. S. *Decision Support Systems: An Organizational Perspective*. Reading, MA: Addison-Wesley, 1978.

[17] Laudon, K. C.; Laudon, J. P. *Management Information Systems: Managing the Digital Firm*. 16th ed. Harlow: Pearson, 2020.

[18] March, J. G.; Simon, H. A. *Organizations*. New York: Wiley, 1958.

[19] Marr, B. *Data Strategy: How to Profit from a World of Big Data, Analytics and AI*. London: Kogan Page, 2017.

[20] McAfee, A.; Brynjolfsson, E. Big data: The management revolution. *Harvard Business Review*, 90(10), 2012, pp. 60-68.

[21] NIST SP 800-53 Rev. 5. *Security and Privacy Controls for Information Systems and Organizations*. Gaithersburg, MD: National Institute of Standards and Technology, 2020.

[22] NIST SP 800-61 Rev. 2. *Computer Security Incident Handling Guide*. Gaithersburg, MD: National Institute of Standards and Technology, 2012.

[23] Power, D. J. *A Brief History of Decision Support Systems*. DSSResources.COM, version 4.0, 2007.

[24] Rothwell, W. J. *Effective Succession Planning: Ensuring Leadership Continuity and Building Talent from Within*. 4th ed. New York: AMACOM, 2010.

[25] Simon, H. A. The New Science of Management Decision. Revised ed. Englewood Cliffs, NJ: Prentice Hall, 1977.

[26] Stone, D. L.; Deadrick, D. L.; Lukaszewski, K. M.; Johnson, R. The influence of technology on the future of human resource management. Human Resource Management Review, 25(2), 2015, pp. 216-231.

[27] Tarafdar, M.; Beath, C. M.; Ross, J. W. Enterprise cognitive computing applications: Opportunities and challenges. IT Professional, 19(4), 2017, pp. 21-27.

ATTACHMENTS

Appendix A: Input Data Schema and Validation Rules

This appendix defines the project input schema and validation constraints used by the analysis engine. It specifies required fields, accepted data types, limits, default values, and units so that manual and CSV-based submissions remain consistent and machine-processable.

Required Fields (14 inputs):

Field	Type	Min	Max	Default	Unit
company_name	Text	-	100 chars	"Company"	-
industry	Category	-	dropdown	"Technology"	-
stage	Category	-	dropdown	"Growth"	-
employee_count	Numeric	1	10000	100	employees
leadership_years	Numeric	0	50	5	years
digital_score	Numeric	1	10	5	1-10 scale
retention_pct	Numeric	0	100	85	%
churn_pct	Numeric	0	100	15	%
dte_ratio	Numeric	0	5	0.5	debt/equity
doc_score	Numeric	0	100	60	% documentation
dep_score	Numeric	0	100	40	% dependency concentration
margin_pct	Numeric	-50	100	15	% profit margin
growth_pct	Numeric	-50	100	20	% YoY growth
cash_months	Numeric	0	60	12	months of runway

CSV Upload Schema (Exact Column Order Required):

The batch parser expects a strict column sequence. If column order is changed, ingestion may fail or map values incorrectly.

```
company_name,industry,stage,employee_count,leadership_years,digital_score,
retention_pct,churn_pct,dte_ratio,doc_score,dep_score,margin_pct,growth_p
ct,cash_months
```

Appendix B: KPI Calculation Formulas and Thresholds

Normalization Functions:

All numeric inputs are normalized to a 0-100 range using domain-specific transformation rules:

- **Leadership Years:** $f(x) = \min(100, 50 + 2 \times x)$ for $x \leq 15$; $\min(100, 80)$ for $x > 15$ (diminishing return after 15 years)
- **Digital Score:** $f(x) = 10 \times x$ (direct linear mapping from 1-10 scale)
- **Retention:** $f(x) = x$ (already 0-100 scale)
- **Churn:** $f(x) = 100 - x$ (inverse: higher churn = higher risk)
- **Debt Pressure:** $f(x) = \min(100, 20 + 16 \times x)$ (scales linearly up to 1:1 ratio)
- **Process Documentation:** $f(x) = x$ (already 0-100 scale)
- **Dependency Concentration:** $f(x) = x$ (already 0-100 scale)
- **Profit Margin:** $f(x) = \max(0, \min(100, 50 + 2 \times x))$
- **Growth Rate:** $f(x) = \max(0, \min(100, 50 + x))$
- **Cash Position:** $f(x) = \max(0, \min(100, 5 \times x))$

Risk Band Thresholds:

Metric	Low Risk	Moderate Risk	High Risk
OHI	70-100	50-69	0-49
LRS	65-100	45-64	0-44
SRS	0-39	40-59	60-100
FSC	65-100	45-64	0-44

Appendix C: API Route Reference and Response Contracts

Authentication Routes:

- ``POST /login`` – Username and password submission; redirects to dashboard on success; returns login form with error message on failure
- ``GET /logout`` – Clears session; redirects to login

Analysis Routes:

- `GET /analysis` – Displays input form with last-used values pre-populated
- `POST /analysis` – Processes form submission; action parameter (run-analysis, import-provider, csv-upload) determines processing path
- `GET /history` – Lists prior analyses with user/role filtering
- `GET /analysis/<id>` – Displays stored analysis details
- `GET /analysis/<id>/pdf` – Streams PDF export as application/pdf

Dashboard Routes:

- `GET /dashboard/<role>` – Renders role-specific dashboard (CEO, HR, Finance, Operations)

Response Context Objects:

AnalysisResult:

```
{
  "company_name": string,
  "analysis_id": integer,
  "created_at": ISO 8601 timestamp,
  "inputs": { 14 input fields },
  "kpis": {
    "lrs": { "score": 0-100, "band": "Low/Moderate/High", "narrative": string },
    "srs": { "score": 0-100, "band": "Low/Moderate/High", "narrative": string },
    "fsc": { "score": 0-100, "band": "Low/Moderate/High", "narrative": string },
    "ohi": { "score": 0-100, "band": "Low/Moderate/High", "narrative": string }
  },
  "ml_prediction": { "probability": 0-100, "confidence": 0-100 },
  "recommendations": [ string, string, ... ],
  "trend": { "prior_ohi": 0-100, "ohi_change": -100 to +100, "direction": "improving/stable/degrading" }
}
```

Appendix D: System Deployment and Operations

Local Development Deployment:

```
# Create and activate virtual environment
python -m venv .venv
```

```

source .venv/bin/activate # On Windows: .venv\Scripts\activate

# Install dependencies
pip install -r requirements.txt

# Initialize database
python -c "from data_store import init_db; init_db()"

# Start Flask server
python app.py

# Access at http://127.0.0.1:5000

```

Demo Credentials:

Username	Password	Role	Purpose
admin	admin123	admin	System administration
ceo_demo	demo123	ceo	CEO dashboard access
hr_demo	demo123	hr	HR dashboard access
finance_demo	demo123	finance	Finance dashboard access
ops_demo	demo123	operations	Operations dashboard access

Production Considerations:

- Replace SQLite with PostgreSQL for concurrent multi-user workloads
- Implement CSRF protection on state-changing forms
- Enable TLS/HTTPS termination and request logging
- Rotate demo credentials; implement LDAP or OAuth integration
- Add structured audit logging for authentication and sensitive queries
- Document data retention and deletion policies
- Implement rate limiting and brute-force protection

Appendix E: Machine Learning Model Details

Training Data Generation:

1. Generate 500 synthetic organizational profiles by random sampling from valid input ranges
2. For each profile, compute deterministic KPIs (LRS, SRS, FSC, OHI)
3. Normalize outputs to probability scale ($OHI/100 = \text{probability estimate}$)
4. Create feature vectors [inputs, normalized_leadership, normalized_digital, ..., normalized_growth]

Model Architecture:

- Algorithm: scikit-learn GradientBoostingRegressor
- Parameters: $n_estimators=100$, $learning_rate=0.1$, $max_depth=3$
- Train/Test split: 80/20 using fixed random seed for reproducibility
- Loss function: Mean Squared Error

Model Performance Metrics:

- Train R^2 : ~ 0.87
- Test R^2 : ~ 0.84
- Mean Absolute Error (MAE): ~ 8.5 (on 0-100 scale)
- Cross-validation (5-fold): $CV\ R^2 = 0.83 \pm 0.03$

Feature Importance Ranking (Approximate):

1. Profit Margin (15%)
2. Digital Score (14%)
3. Growth Rate (12%)
4. Retention Percentage (11%)
5. Churn Percentage (10%)
6. Cash Position (9%)
7. Leadership Years (8%)
8. Debt-to-Equity Ratio (7%)
9. Process Documentation (7%)

10. Key-Person Dependency (5%)

Limitations:

- Model trained on synthetic data, not historical organizational outcomes
- No calibration against real performance data
- Feature importance approximation is directional, not exact (not SHAP values)
- Model assumes linear separability in feature space; nonlinear patterns may be missed

Appendix F: Validation Test Cases

Test Case 1: Input Boundary Validation

Input: leadership_years = -5 (invalid, out of range) Expected: Clamped to 0; validation note "leadership_years clamped to minimum (0)" Observed: ✓ Correct

Test Case 2: Missing Value Substitution

Input: profit_margin = null Expected: Substituted with default (15%); note "margin_pct defaulted to 15" Observed: ✓ Correct

Test Case 3: Deterministic Computation

Input: Same company profile submitted twice Expected: Byte-identical output Observed: ✓ Correct

Test Case 4: Role-Based Access Control

Action: Non-admin user attempts to view another user's analysis Expected: Access denied; redirect to history Observed: ✓ Correct

Test Case 5: CSV Batch Robustness

Input: CSV with 50 rows; 5 rows containing invalid industry codes Expected: Valid rows processed; invalid rows marked with notes; no route failure Observed: ✓ Correct; 45 analyses persisted, 5 noted with validation issues

Test Case 6: PDF Export Reproducibility

Action: Export same analysis twice Expected: Byte-identical PDFs Observed: ✓
Correct

Appendix G: End-to-End Workflow Specifications by User Role

This appendix documents the exact operational workflow executed by each role. The objective is to ensure that the implementation can be audited not only at code level, but also at process level. Each workflow defines entry conditions, sequence steps, expected system responses, and outcome artifacts.

G.1 CEO Workflow (Strategic Readiness Review)

Entry Conditions:

- Authenticated role = `ceo` or `admin`
- At least one company analysis record exists in storage
- Dashboard cache includes last KPI snapshot or can be computed dynamically

Sequence Steps:

1. User opens `/dashboard/ceo`
2. System loads latest analysis for user scope (global if admin)
3. KPI panel renders OHI, LRS, SRS, FSC with color-coded risk bands
4. Strategic recommendation block displays top-priority interventions
5. Trend panel compares current OHI against previous snapshot
6. User opens history table and selects company for deep dive
7. User exports selected analysis PDF for board reporting

Expected Outputs:

- Strategic view with integrated organizational health
- Ranked action priorities with explanatory narratives
- Audit-ready PDF summary for management communication

Control Points:

- Role authorization on route access
- Deterministic KPI rendering for reproducibility
- Snapshot consistency between dashboard and export

G.2 HR Workflow (Retention and Leadership Stability Review)**Entry Conditions:**

- Authenticated role = `hr` or `admin`
- Input contains retention/churn/leadership indicators

Sequence Steps:

1. User opens `/analysis` and enters HR-focused values
2. User submits with action `run-analysis`
3. Engine validates fields and applies boundary clamping if needed
4. LRS and SRS subcomponents compute with transparent formulas
5. HR dashboard renders attrition-sensitive risk narratives
6. User compares trend against prior submissions
7. User records intervention plan from recommendations

Expected Outputs:

- Leadership readiness interpretation aligned to role scope
- Retention-driven risk exposure indicators
- Action list suitable for people operations planning

G.3 Finance Workflow (Stability and Liquidity Review)**Entry Conditions:**

- Authenticated role = `finance` or `admin`

- Financial indicators available (margin, growth, debt, cash)

Sequence Steps:

1. User opens finance dashboard route
2. System computes/loads FSC and OHI components
3. Debt pressure and cash runway contributions are visualized
4. User reviews leverage-driven recommendation text
5. User exports report for quarterly governance package

Expected Outputs:

- Financial risk interpretation in normalized 0-100 scale
- Explicit documentation of metric-to-score mapping
- Reproducible artifact for finance committee review

G.4 Operations Workflow (Process and Dependency Risk Review)

Entry Conditions:

- Authenticated role = `operations` or `admin`
- Documentation score and dependency concentration supplied

Sequence Steps:

1. User accesses operations dashboard
2. System presents process maturity and concentration indicators
3. SRS contribution highlights bottlenecks and key-person exposure
4. Recommendations emphasize cross-training and documentation actions
5. User tracks post-intervention trend in subsequent runs

Expected Outputs:

- Operational risk visibility with defensible formula traceability
- Prioritized mitigation guidance for resilience improvement

- Historical evidence of intervention effectiveness

G.5 Batch Analyst Workflow (Multi-Company Screening)

Entry Conditions:

- User has valid CSV in required schema order
- Role authorized for batch analysis

Sequence Steps:

1. User uploads CSV through `/analysis` action `csv-upload`
2. Parser validates row-level format and value ranges
3. Each row processed through deterministic engine
4. Validation notes attached for rows with corrections/defaults
5. Batch ranking generated by OHI/SRS logic
6. Results persisted and exposed in history page

Expected Outputs:

- Ranked multi-company overview for rapid triage
- Transparent notes for data quality anomalies
- Persistent evidence for repeated comparative analysis

Appendix H: Comprehensive Metric Dictionary and Interpretation Rules

This appendix defines each primary input and derived score in implementation language suitable for auditors, maintainers, and future research teams.

H.1 Input Field Definitions

1. company_name

Purpose: Human-readable company identifier in UI, reports, and history. Validation: Non-empty string, maximum length 100. Operational Note: Used as grouping key for trend narratives.

2. **industry**

Purpose: Domain context for interpretation and comparison. Validation: Enum against configured industry list. Operational Note: Unknown value handled by fallback category + note.

3. **stage**

Purpose: Lifecycle classification (e.g., startup, growth, mature). Validation: Enum against configured stages. Operational Note: Influences recommendation language, not core formula weights.

4. **employee_count**

Purpose: Organization scale context for dashboard summaries. Validation: Integer range 1-10000. Operational Note: Primarily descriptive in current model version.

5. **leadership_years**

Purpose: Leadership maturity proxy. Validation: Numeric range 0-50. Transformation: Normalized with diminishing return threshold around 15 years.

6. **digital_score**

Purpose: Digital capability readiness indicator. Validation: Numeric range 1-10. Transformation: Linear map to 0-100 scale.

7. **retention_pct**

Purpose: Workforce stability measure. Validation: Numeric 0-100. Transformation: Direct mapping.

8. **churn_pct**

Purpose: Attrition pressure measure. Validation: Numeric 0-100. Transformation: Inverse mapping to represent risk escalation.

9. **dte_ratio**

Purpose: Leverage exposure indicator. Validation: Numeric 0-5. Transformation: Linear pressure mapping with upper cap.

10. **doc_score**

Purpose: Process documentation maturity. Validation: Numeric 0-100.
Transformation: Direct mapping.

11. **dep_score**

Purpose: Key-person dependency concentration. Validation: Numeric 0-100.
Transformation: Direct mapping into scaling risk contribution.

12. **margin_pct**

Purpose: Profitability resilience. Validation: Numeric -50 to 100. Transformation:
Shifted linear mapping to 0-100 interval.

13. **growth_pct**

Purpose: Expansion momentum. Validation: Numeric -50 to 100. Transformation:
Shifted linear mapping with clipping.

14. **cash_months**

Purpose: Liquidity runway. Validation: Numeric 0-60. Transformation: Scaled to
0-100 with upper cap.

H.2 Derived KPI Definitions

1. LRS (Leadership Readiness Score)

Interpretation: Ability of leadership and people systems to sustain change. Expected
Usage: HR and CEO readiness reviews.

2. SRS (Scaling Risk Score)

Interpretation: Probability of operational stress under growth conditions. Expected
Usage: Operations and HR risk monitoring.

3. FSC (Financial Stability Composite)

Interpretation: Resilience of margin, growth, leverage, and liquidity profile.
Expected Usage: Finance governance and quarterly monitoring.

4. OHI (Organizational Health Index)

Interpretation: Integrated holistic indicator across all principal dimensions. Expected
Usage: Cross-functional strategic decision support.